

M507_Individual_Final_project(Jupyter Notebook_code)_GH1019253

June 30, 2022

1 Final Assessment for M507(Methods of Prediction) Group c

```
[1]: from IPython.display import Image  
Image("GISMA_LOGO.png",width = 200, height = 200)
```

```
[1]:
```

GISMA

BUSINESS
SCHOOL

1.0.1 Predicting Anamoly in Credit Card Transactions

1.1 Table of the Content

1.1.1 1) Introduction

1.1.2 2) Problem Statement

1.1.3 3) Methodology & Apporach

1.1.4 4) Importing Libraries

1.1.5 5) Importing Dataset

1.1.6 6) Data Exploration

6.1) Checking for Missing Values

6.2) Data Visualization

1.1.7 7) Data Preprocessing

7.1) Scaling the Data using Standard Scaler

1.1.8 8) Train/Test Split

1.1.9 9) Model Building

9.1)Random Forest Classifier

9.2) Decision Tree Classifier

9.3) Gradient Boosting Classifier

9.4) SGDClassifier

1.1.10 10) Neural Network Models

10.1)Neural Network using Weight Class

10.2) Neural Network using Undersampling Technique

10.3) Neural Network using SMOTE 10.3.1) Testing on Full Data Set Now

1.1.11 11) Model Evaluation

1.1.12 12) Conclusion and Actionable Insights

2 1) Introduction

I have been hired as a Data Scientist in a Europeon Central bank and my duty is to provide end-to-end solutions related to Data. Nowadays we have seen many kinds of scams in the banking sector and one of the major scam we have noticed is in the Credit Card Transactions.The data that has

been provided to me is highly confidential and it contains only numerical values which are result of PCA Transformation. The features v1 ,v2 v28 are the result of PCA Transformation and Time and Amount column is in original condition to use in prediction. The data we have obtained is based on 284,807 transactions out of which 492 transactions are Fraudulent and these transaction are recorded in September 2013 and these transaction were made by European card holders. As a data scientist now that's my job to make a system and can detect these kind of Anomalies.

- DataSet link (<https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud>)

3 2) Problem Statement

As I have discussed above we have the data of credit card transactions which are about 284,807 transactions and around 492 are Fraudulent, so we can see that the data is highly Imbalanced and we have to resolve this issue and I have to explore data in detail to find out the correlations in data and to clear out the Outliers in data. Due to Imbalancing we can not fit data in neural networks so I have made it balanced by using advance balancing techniques and we have to make a system which can detect these anomaly in future with high precision and for that I will use several Algorithms to find out which will work best in this case.

The Major Business question in this task is that: - 1) How to deal with Imbalance data ? - 2) Which approach will work better in this condition? - 3) What algorithm we should use in this kind of problems?

4 3) Methodology and Approach

First thing to notice here is that the data is of Numerical nature and it's seems highly Imbalanced, So First of all I have to do Exploratory Data Analysis to find out the hidden knowledge in data, later on after data findings I will move forward towards data preprocessing and cleansing if it will be needed. Once my data will be cleaned I will have to scale it and later on I have to make the data balance to feed into the Machine Learning models. Once I will get the results for different model then I will do model evaluation and then I can concluded what can be done in future to overcome such problems and how we can improve more.

5 4) Importing Libraries

```
[2]: import pandas as pd
import numpy as np

# import keras
# from keras.models import Sequential
# from keras.layers import Dense
# from keras.layers import Dropout

import keras
from keras import layers
from keras.models import Sequential
from keras.layers.core import Dense, Dropout
```

```

from sklearn.utils import class_weight
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier

from imblearn.over_sampling import SMOTE

import matplotlib.pyplot as plt
import itertools

import sklearn.model_selection
import sklearn.metrics

import sklearn.neighbors

from sklearn import svm, datasets
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score, plot_roc_curve, classification_report
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import GradientBoostingClassifier
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns
color = sns.color_palette()
%matplotlib inline

import missingno as msno
import warnings
warnings.filterwarnings('ignore')

sns.set(style='whitegrid', palette='muted', font_scale=1.5)

from pylab import rcParams
rcParams['figure.figsize'] = 14, 8

```

6 5) Importing Dataset

```

[3]: CC_data = pd.read_csv('data/creditcard.csv')

CC_data.head()

```

```
[3]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	\
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	

	V8	V9	...	V21	V22	V23	V24	V25	\
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539	
1	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170	
2	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.327642	
3	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376	
4	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010	

	V26	V27	V28	Amount	Class
0	-0.189115	0.133558	-0.021053	149.62	0
1	0.125895	-0.008983	0.014724	2.69	0
2	-0.139097	-0.055353	-0.059752	378.66	0
3	-0.221929	0.062723	0.061458	123.50	0
4	0.502292	0.219422	0.215153	69.99	0

[5 rows x 31 columns]

7 6) Data Exploration

```
[4]: CC_data.columns
```

```
[4]: Index(['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10',
          'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20',
          'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount',
          'Class'],
          dtype='object')
```

- Here we can that we have 31 columns and in this dataset which contains 284,807 transactions(rows)

```
[5]: CC_data.shape
```

```
[5]: (284807, 31)
```

```
[6]: CC_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column    Non-Null Count  Dtype
---  -

```

```

0   Time      284807 non-null float64
1   V1        284807 non-null float64
2   V2        284807 non-null float64
3   V3        284807 non-null float64
4   V4        284807 non-null float64
5   V5        284807 non-null float64
6   V6        284807 non-null float64
7   V7        284807 non-null float64
8   V8        284807 non-null float64
9   V9        284807 non-null float64
10  V10       284807 non-null float64
11  V11       284807 non-null float64
12  V12       284807 non-null float64
13  V13       284807 non-null float64
14  V14       284807 non-null float64
15  V15       284807 non-null float64
16  V16       284807 non-null float64
17  V17       284807 non-null float64
18  V18       284807 non-null float64
19  V19       284807 non-null float64
20  V20       284807 non-null float64
21  V21       284807 non-null float64
22  V22       284807 non-null float64
23  V23       284807 non-null float64
24  V24       284807 non-null float64
25  V25       284807 non-null float64
26  V26       284807 non-null float64
27  V27       284807 non-null float64
28  V28       284807 non-null float64
29  Amount    284807 non-null float64
30  Class     284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB

```

7.1 6.1) Checking for Missing Values

```
[7]: CC_data.isnull().sum()
```

```

[7]: Time      0
     V1        0
     V2        0
     V3        0
     V4        0
     V5        0
     V6        0
     V7        0
     V8        0

```


```
V9          0
V10         0
V11         0
V12         0
V13         0
V14         0
V15         0
V16         0
V17         0
V18         0
V19         0
V20         0
V21         0
V22         0
V23         0
V24         0
V25         0
V26         0
V27         0
V28         0
Amount      0
Class       0
dtype: int64
```

```
[8]: print(CC_data.shape)
      print(CC_data["Class"].value_counts())
```

```
(284807, 31)
0    284315
1      492
Name: Class, dtype: int64
```

```
[9]: msno.bar(CC_data)
```

```
[9]: <AxesSubplot:>
```

- By Plotting the values of the data frame we can see there is no null values in the data.

7.2 6.2) Data Visualization


```
[10]: fig, axes = plt.subplots(ncols=2, figsize=(17, 5), dpi=100)
plt.tight_layout()

CC_data["Class"].value_counts().plot(kind='pie', ax=axes[0], labels=['Normal_
↳ Transaction (284315)', 'Fradulent Transaction (492)'])
temp = CC_data["Class"].value_counts()
sns.barplot(temp.index, temp, ax=axes[1])

axes[0].set_ylabel(' ')
axes[1].set_ylabel(' ')
axes[1].set_xticklabels(["Normal Transaction (284315)", "Fradulent Transaction_
↳ (492)"])

axes[0].set_title('Target Distribution in Dataset', fontsize=13)
axes[1].set_title('Target Count in Dataset', fontsize=13)

plt.show()
```



- Data visualization states that there are 284,315 transactions that are normal and 492 are of Fraudulent type, It's clear that data is highly Imbalanced and we have to take care of that in future before feeding the data in ML Models

```
[11]: fraud_trans = CC_data[CC_data.Class == 1]
      normal_trans = CC_data[CC_data.Class == 0]
```

```
[12]: fraud_trans.shape
```

```
[12]: (492, 31)
```

```
[13]: normal_trans.shape
```

```
[13]: (284315, 31)
```

```
[14]: fraud_trans.Amount.describe()
```

```
[14]: count      492.000000
      mean       122.211321
      std       256.683288
      min        0.000000
      25%        1.000000
      50%        9.250000
      75%       105.890000
      max       2125.870000
      Name: Amount, dtype: float64
```

```
[15]: normal_trans.Amount.describe()
```

```
[15]: count      284315.000000
      mean        88.291022
      std       250.105092
      min        0.000000
      25%        5.650000
      50%       22.000000
      75%       77.050000
```

```
max      25691.160000
Name: Amount, dtype: float64
```


```
[16]: f, (ax1, ax2) = plt.subplots(2, 1, sharex=True)
f.suptitle('Amount per transaction by class')

bins = 50

ax1.hist(fraud_trans.Amount, bins = bins)
ax1.set_title('Fraud')

ax2.hist(normal_trans.Amount, bins = bins)
ax2.set_title('Normal')

plt.xlabel('Amount ($)')
plt.ylabel('Number of Transactions')
plt.xlim((0, 20000))
plt.yscale('log')
plt.show();
```




- Here we can observe that in most of the fraud transactions the amount spent is around 300 dollars and in Normal transactions it can vary between 1 dollar to 10,000 and there are some outliers as well.

```
[17]: f, (ax1, ax2) = plt.subplots(2, 1, sharex=True)
f.suptitle('Time of transaction vs Amount by class')

ax1.scatter(fraud_trans.Time, fraud_trans.Amount)
ax1.set_title('Fraud')

ax2.scatter(normal_trans.Time, normal_trans.Amount)
ax2.set_title('Normal')

plt.xlabel('Time (in Seconds)')
plt.ylabel('Amount')
plt.show()
```



- As we can see that there are numbers of transactions which were held at same time so it's hard to state that what time the fraudulent transactions were made, although we can see that there are some outliers in both so we have figure them out.
- One thing to note here is that time doesn't really matter to detect the fraudulent transactions.

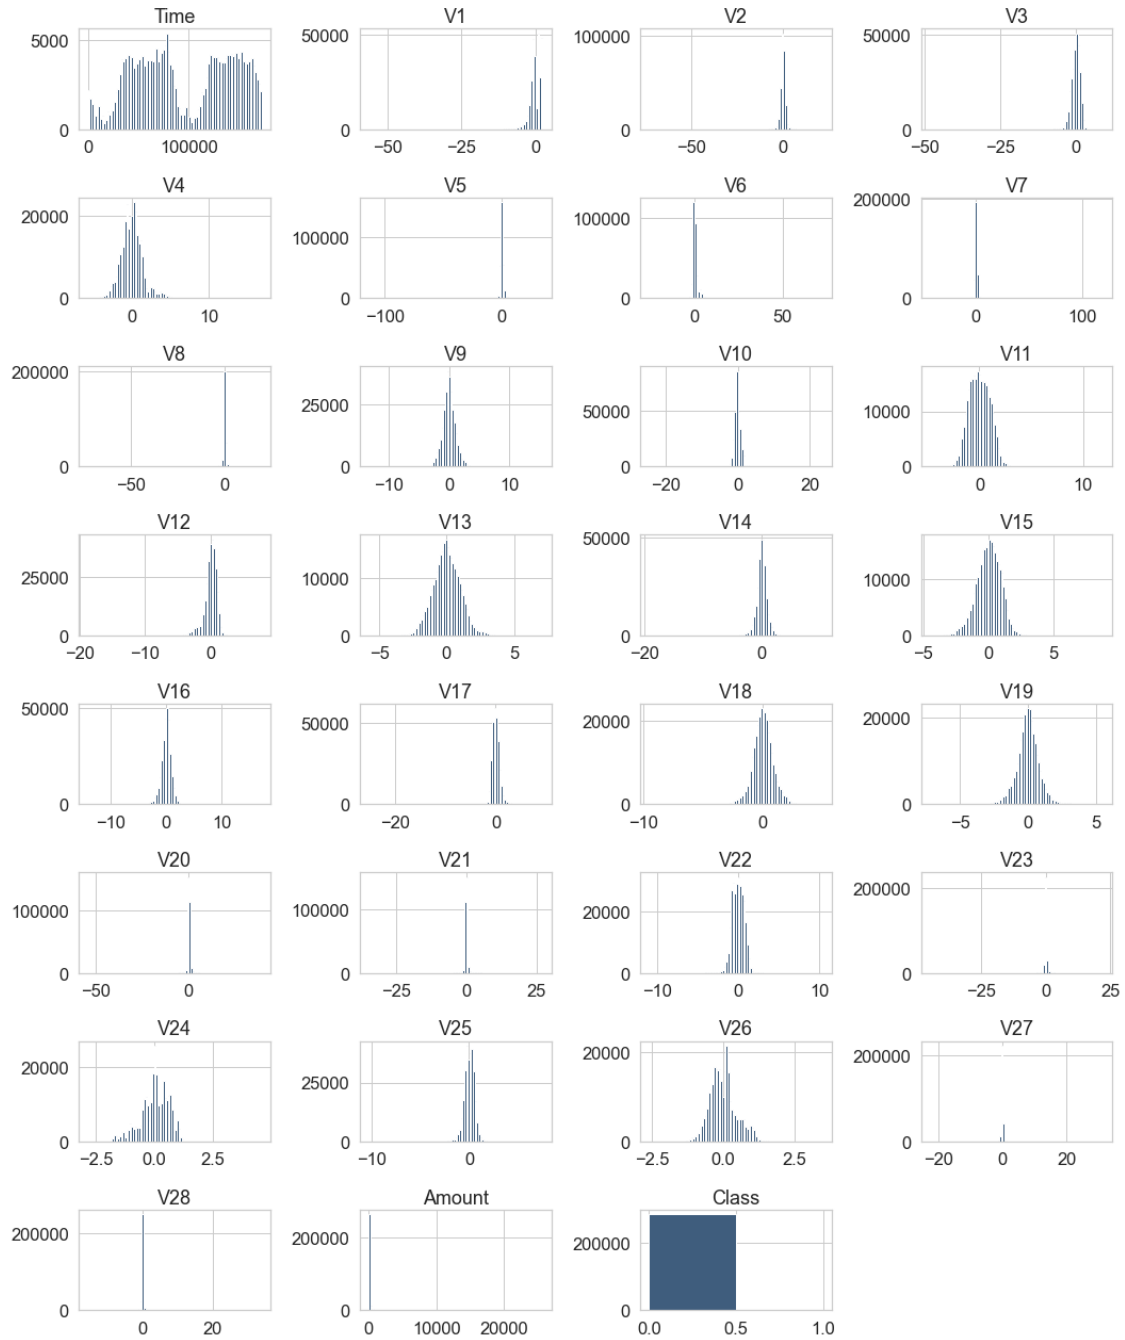
```
[18]: ## Histograms
fig = plt.figure(figsize=(15, 20))
plt.suptitle('Plotting Histograms of Features', fontsize=20)
for i in range(CC_data.shape[1]):
    plt.subplot(8, 4, i + 1)
    f = plt.gca()
```

```
f.set_title(CC_data.columns.values[i])

vals = np.size(CC_data.iloc[:, i].unique())
if vals >= 100:
    vals = 100


plt.hist(CC_data.iloc[:, i], bins=vals, color='#3F5D7D')
plt.tight_layout(rect=[0, 0.03, 1, 0.95])
```

Plotting Histograms of Features



- Here we can observe that most of the data is of Class 0 (Non-Fraudulent) and very Minor data is right Skewed.

```
[19]: ## Linear Correlation with Response Variable (Note: Models like RandomForest
      ↪are not linear)
      CC_df = CC_data.drop(columns = ['Class'])    # drop non numerical columns
      CC_df.corrwith(CC_data.Class).plot.bar(
          figsize = (20, 10), title = "Correlation with Class Fraudulent or Not",
          ↪fontsize = 15,
          rot = 45, grid = True)
      plt.show()
```



8 7) Data Preprocessing

8.1 7.1) Scaling the Data using Standard Scaler

```
[20]: CC_data['Amount_normalized'] = StandardScaler().fit_transform(CC_data['Amount']).
      ↪values.reshape(-1,1))
      CC_data = CC_data.drop(['Amount'],axis=1)
```

- I have used Standard Scaler here to normalize the Amount column after finding that we need to normalize it as every other column is already normalized.

```
[21]: CC_data.head()
```

```
[21]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	\
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	

```
4 2.0 -1.158233 0.877737 1.548718 0.403034 -0.407193 0.095921 0.592941
```

```

      V8      V9  ...      V21      V22      V23      V24      V25  \
0 0.098698 0.363787 ... -0.018307 0.277838 -0.110474 0.066928 0.128539
1 0.085102 -0.255425 ... -0.225775 -0.638672 0.101288 -0.339846 0.167170
2 0.247676 -1.514654 ... 0.247998 0.771679 0.909412 -0.689281 -0.327642
3 0.377436 -1.387024 ... -0.108300 0.005274 -0.190321 -1.175575 0.647376
4 -0.270533 0.817739 ... -0.009431 0.798278 -0.137458 0.141267 -0.206010

```

```

      V26      V27      V28  Class  Amount_normalized
0 -0.189115 0.133558 -0.021053      0      0.244964
1 0.125895 -0.008983 0.014724      0     -0.342475
2 -0.139097 -0.055353 -0.059752      0      1.160686
3 -0.221929 0.062723 0.061458      0      0.140534
4 0.502292 0.219422 0.215153      0     -0.073403

```

[5 rows x 31 columns]

```
[22]: CC_data = CC_data.drop(['Time'],axis=1)
      CC_data.head()
```

```
[22]:
      V1      V2      V3      V4      V5      V6      V7  \
0 -1.359807 -0.072781 2.536347 1.378155 -0.338321 0.462388 0.239599
1 1.191857 0.266151 0.166480 0.448154 0.060018 -0.082361 -0.078803
2 -1.358354 -1.340163 1.773209 0.379780 -0.503198 1.800499 0.791461
3 -0.966272 -0.185226 1.792993 -0.863291 -0.010309 1.247203 0.237609
4 -1.158233 0.877737 1.548718 0.403034 -0.407193 0.095921 0.592941

```

```

      V8      V9      V10  ...      V21      V22      V23      V24  \
0 0.098698 0.363787 0.090794 ... -0.018307 0.277838 -0.110474 0.066928
1 0.085102 -0.255425 -0.166974 ... -0.225775 -0.638672 0.101288 -0.339846
2 0.247676 -1.514654 0.207643 ... 0.247998 0.771679 0.909412 -0.689281
3 0.377436 -1.387024 -0.054952 ... -0.108300 0.005274 -0.190321 -1.175575
4 -0.270533 0.817739 0.753074 ... -0.009431 0.798278 -0.137458 0.141267

```

```

      V25      V26      V27      V28  Class  Amount_normalized
0 0.128539 -0.189115 0.133558 -0.021053      0      0.244964
1 0.167170 0.125895 -0.008983 0.014724      0     -0.342475
2 -0.327642 -0.139097 -0.055353 -0.059752      0      1.160686
3 0.647376 -0.221929 0.062723 0.061458      0      0.140534
4 -0.206010 0.502292 0.219422 0.215153      0     -0.073403

```

[5 rows x 30 columns]

- Here I have dropped the Time column as it is not relevant to our prediction.


```
[23]: X = CC_data.iloc[:, CC_data.columns != 'Class']
      y = CC_data.iloc[:, CC_data.columns == 'Class'] # Response variable
      ↪determining if fraudulent or not
```

```
[24]: y.head()
```

```
[24]:   Class
      0    0
      1    0
      2    0
      3    0
      4    0
```

```
[25]: X.shape
```

```
[25]: (284807, 29)
```

```
[26]: y.shape
```

```
[26]: (284807, 1)
```

- Here we can see that our data is in shape after scaling to feed into Machine Learning Model.

9 8) Train/Test Split

```
[27]: X_train, X_test, y_train, y_test = train_test_split(X,y, test_size = 0.2,
      ↪random_state=10)
```

```
[28]: X_train.shape
```

```
[28]: (227845, 29)
```

```
[29]: X_test.shape
```

```
[29]: (56962, 29)
```

- In the above section I have splitted the data into train, test set , I have used 80 percent for training and 20 percent for Testing.

10 9) Model Building

10.1 9.1)Random Forest Classifier

```
[30]: %%time
      parameters = {
```

```

        'n_estimators':range(10,100,10),
        'criterion':['gini','entropy'],
        'max_leaf_nodes':range(2,10,1),
        'max_features':['auto','log2']

    }

    RF_classifier = RandomForestClassifier()

    Grid_Sr = RandomizedSearchCV(estimator = RF_classifier, param_distributions = {
        'parameters', n_iter = 10, cv = 3,n_jobs = -1)
    Grid_Sr.fit(X_train,y_train.values.ravel())

```

CPU times: total: 29.2 s

Wall time: 2min 6s

```

[30]: RandomizedSearchCV(cv=3, estimator=RandomForestClassifier(), n_jobs=-1,
        param_distributions={'criterion': ['gini', 'entropy'],
                              'max_features': ['auto', 'log2'],
                              'max_leaf_nodes': range(2, 10),
                              'n_estimators': range(10, 100, 10)})

```

```

[31]: Grid_Sr.best_params_

```

```

[31]: {'n_estimators': 40,
        'max_leaf_nodes': 8,
        'max_features': 'auto',
        'criterion': 'entropy'}

```

- Here I have used RandomizedSearchCV to get the best parameters to feed into Random Forest Classifier and as we can see we have got the best parameters which i will be using now

```

[32]: ran_for = RandomForestClassifier(**Grid_Sr.best_params_)

    ran_for.fit(X_train,y_train.values.ravel())

```

```

[32]: RandomForestClassifier(criterion='entropy', max_features='auto',
                              max_leaf_nodes=8, n_estimators=40)

```

```

[33]: y_predict = ran_for.predict(X_test)

```

```

[34]: ran_for.score(X_test,y_test)

```

```

[34]: 0.9995084442259752

```

```

[35]: accuracy = accuracy_score(y_test, y_predict)
    precision = precision_score(y_test, y_predict)

```

```

recall = recall_score(y_test, y_predict)
f1 = f1_score(y_test, y_predict)

print('accuracy:%0.4f'%accuracy, '\tprecision:%0.4f'%precision, '\trecall:%0.
↪4f'%recall, '\tF1-score:%0.4f'%f1)

```


accuracy:0.9995 precision:0.9125 recall:0.7766 F1-score:0.8391

- We can see that we got a very good Train and Test accuracy here using RandomizedSearch for RandomForestClassifier

```

[36]: con_max = confusion_matrix(y_predict, y_test)
sns.heatmap(con_max, annot=True);

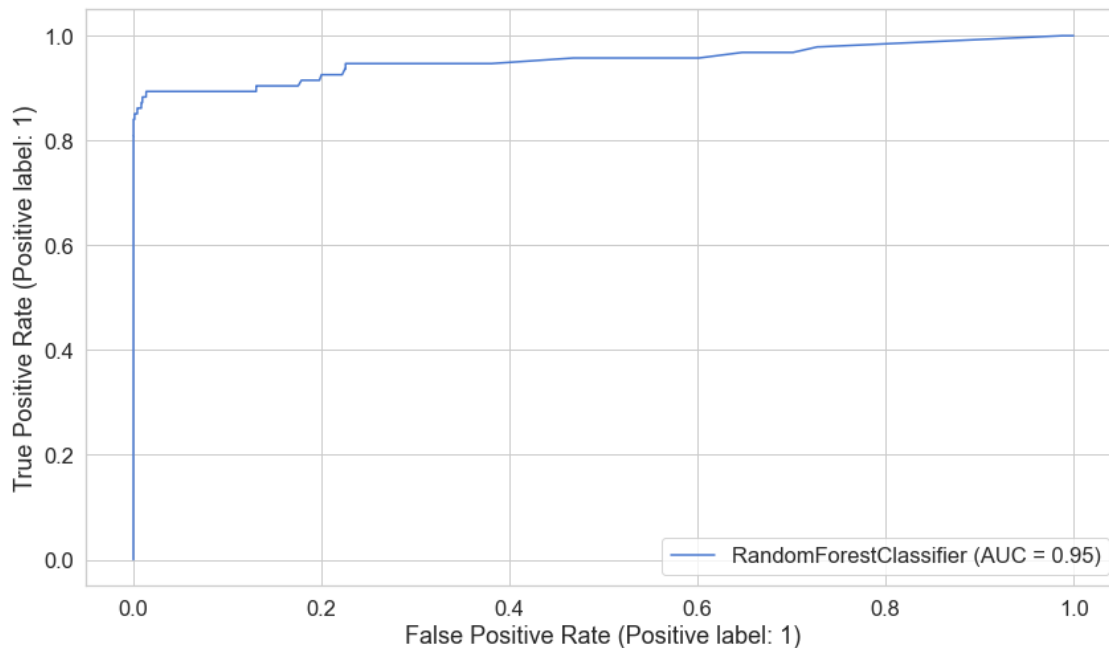
```



```

[37]: ROC_RF = plot_roc_curve(ran_for, X_test, y_test)
plt.show()

```



```
[38]: ### Store results in dataframe for comparing various Models
Testing_result = pd.DataFrame(['RandomForest', accuracy, 1-recall, recall,
    ↪ precision, f1]),
    columns = ['Model', 'Accuracy', 'FalseNegRate', 'Recall',
    ↪ 'Precision', 'F1 Score'])
Testing_result
```

```
[38]:      Model  Accuracy  FalseNegRate  Recall  Precision  F1 Score
0  RandomForest  0.999508      0.223404  0.776596      0.9125   0.83908
```

- I am appending the result of the model to in a data frame to use later for Model Evaluation

10.2 9.2) Decision Tree Classifier

```
[39]: %%time
dec_tree = DecisionTreeClassifier(criterion='gini', random_state=100,
    ↪ max_depth=3, min_samples_leaf=5)
dec_tree.fit(X_train,y_train.values.ravel())
```

```
CPU times: total: 2.92 s
Wall time: 2.93 s
```

```
[39]: DecisionTreeClassifier(max_depth=3, min_samples_leaf=5, random_state=100)
```

- Now I have used Decision Tree Classifier with some random parameters to feed into model (here I haven't used GridSearch or RandomizedSearch as it is time taking process)

```
[40]: y_predicted = dec_tree.predict(X_test)
      dec_tree.score(X_test,y_test)
```


```
[40]: 0.999420666409185
```

```
[41]: accuracy = accuracy_score(y_test, y_predicted)
      precision = precision_score(y_test, y_predicted)
      recall = recall_score(y_test, y_predicted)
      f1 = f1_score(y_test, y_predicted)
      print('accuracy:%0.4f'%accuracy, '\tprecision:%0.4f'%precision, '\trecall:%0.
            ↪4f'%recall, '\tF1-score:%0.4f'%f1)
```

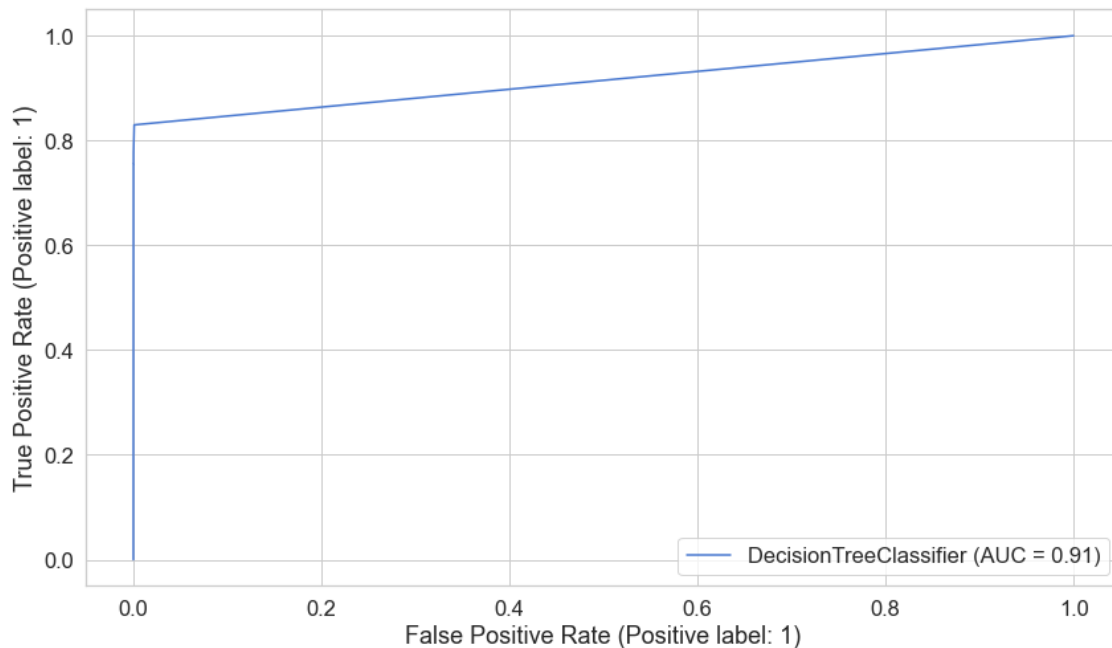
```
accuracy:0.9994      precision:0.8765      recall:0.7553      F1-score:0.8114
```

- Even though I selected the parameters randomly but still we got a very good train/test accuracy which is a good sign.

```
[42]: con_max = confusion_matrix(y_predicted, y_test)
      sns.heatmap(con_max, annot=True);
```



```
[43]: ROC_DT = plot_roc_curve(dec_tree, X_test, y_test)
      plt.show()
```



```
[44]: Mod_result = pd.DataFrame([['DecisionTree', accuracy, 1-recall, recall,
    ↪precision, f1]],
    columns = ['Model', 'Accuracy', 'FalseNegRate', 'Recall',
    ↪'Precision', 'F1 Score'])
Testing_result = Testing_result.append(Mod_result, ignore_index = True)
Testing_result
```

```
[44]:
```

	Model	Accuracy	FalseNegRate	Recall	Precision	F1 Score
0	RandomForest	0.999508	0.223404	0.776596	0.912500	0.839080
1	DecisionTree	0.999421	0.244681	0.755319	0.876543	0.811429

10.3 9.3) Gradient Boosting Classifier

```
[45]: %%time

grad_Boost_params ={
    'n_estimators': 100,
    'max_features': 0.9,
    'max_depth': 4,
    'min_samples_leaf': 2,
    'max_features' : 'sqrt',
    'verbose': 0
}
```

CPU times: total: 0 ns

Wall time: 0 ns

```
[46]: grad_boost = GradientBoostingClassifier(**grad_Boost_params)
      # GS.fit(X_train,y_train.values.ravel())
      grad_boost.fit(X_train,y_train.values.ravel())
```

```
[46]: GradientBoostingClassifier(max_depth=4, max_features='sqrt', min_samples_leaf=2)
```

```
[47]: y_predicted = grad_boost.predict(X_test)
```

```
[48]: grad_boost.score(X_test,y_test)
```

```
[48]: 0.9994382219725431
```


```
[49]: accuracy = accuracy_score(y_test, y_predicted)
      precision = precision_score(y_test, y_predicted)
      recall = recall_score(y_test, y_predicted)
      f1 = f1_score(y_test, y_predicted)

      print('accuracy:%0.4f'%accuracy, '\tprecision:%0.4f'%precision, '\trecall:%0.
      ↪4f'%recall, '\tF1-score:%0.4f'%f1)
```

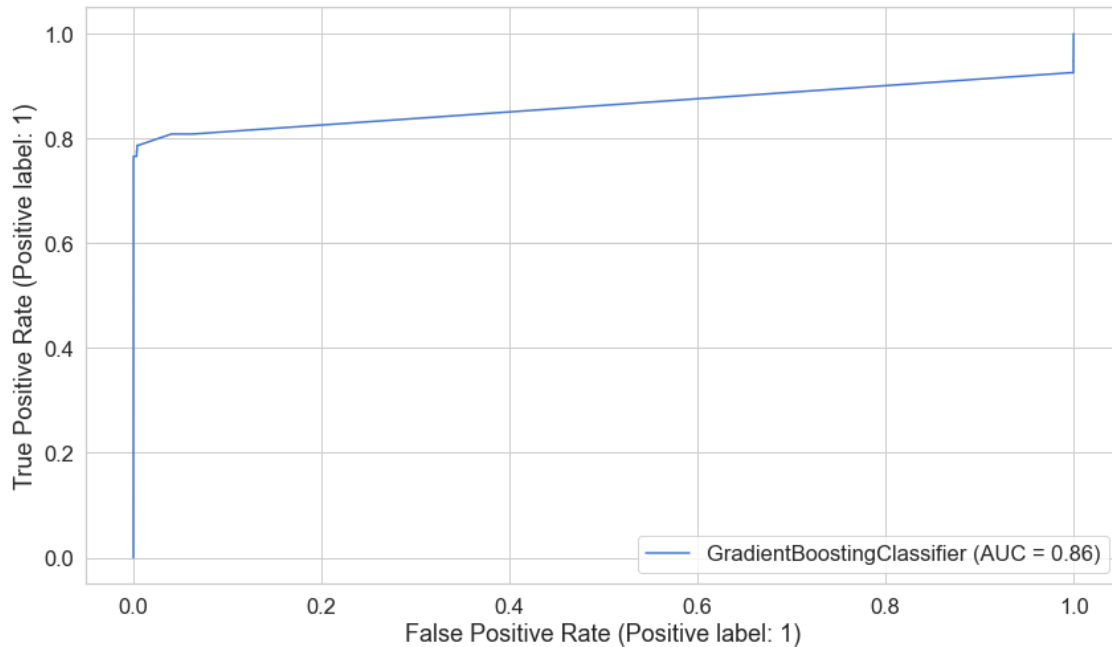
accuracy:0.9994 precision:0.8780 recall:0.7660 F1-score:0.8182

- I have defined Gradient Boosting Classifier with grid of parameters to select and we have seen very good result which is almost 100 percent

```
[50]: con_max = confusion_matrix(y_predicted, y_test)
      sns.heatmap(con_max, annot=True);
```



```
[51]: ROC_RF = plot_roc_curve(grad_boost, X_test, y_test)
plt.show()
```



```
[52]: Mod_result = pd.DataFrame(['Grad_Boost', accuracy, 1-recall, recall,
    ↪ precision, f1]),
    columns = ['Model', 'Accuracy', 'FalseNegRate', 'Recall',
    ↪ 'Precision', 'F1 Score'])
Testing_result = Testing_result.append(Mod_result, ignore_index = True)
Testing_result
```

```
[52]:
```

	Model	Accuracy	FalseNegRate	Recall	Precision	F1 Score
0	RandomForest	0.999508	0.223404	0.776596	0.912500	0.839080
1	DecisionTree	0.999421	0.244681	0.755319	0.876543	0.811429
2	Grad_Boost	0.999438	0.234043	0.765957	0.878049	0.818182

10.4 9.4) SGDClassifier

```
[53]: %%time
from sklearn.linear_model import SGDClassifier
sgd_clf = SGDClassifier(penalty="l2", random_state=0, max_iter=1000)
sgd_clf.fit(X_train,y_train.values.ravel())
```

CPU times: total: 562 ms

Wall time: 560 ms

```
[53]: SGDClassifier(random_state=0)
```

```
[54]: y_pred = grad_boost.predict(X_test)
```

```
[55]: sgd_clf.score(X_test,y_test)
```

```
[55]: 0.9992099996488887
```


```
[56]: accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print('accuracy:%0.4f'%accuracy, '\tprecision:%0.4f'%precision, '\trecall:%0.
↪4f'%recall, '\tF1-score:%0.4f'%f1)
```

accuracy:0.9994 precision:0.8780 recall:0.7660 F1-score:0.8182

- Here I have used Stochastic Gradient Descent (SGD) Classifier and the is nearly same as Gradient Boosting Classifier

```
[57]: con_max = confusion_matrix(y_pred, y_test)
sns.heatmap(con_max, annot=True);
```



```
[58]: Mod_result = pd.DataFrame([['SGD_Classifier', accuracy, 1-recall, recall,
    ↪precision, f1]],
    columns = ['Model', 'Accuracy', 'FalseNegRate', 'Recall',
    ↪'Precision', 'F1 Score'])
Testing_result = Testing_result.append(Mod_result, ignore_index = True)
Testing_result
```

```
[58]:
```

	Model	Accuracy	FalseNegRate	Recall	Precision	F1 Score
0	RandomForest	0.999508	0.223404	0.776596	0.912500	0.839080
1	DecisionTree	0.999421	0.244681	0.755319	0.876543	0.811429
2	Grad_Boost	0.999438	0.234043	0.765957	0.878049	0.818182
3	SGD_Classifier	0.999438	0.234043	0.765957	0.878049	0.818182

11 10) Neural Network Models

Now on I will be working with Neural Networks I have defined the model using Keras Sequential API with Dense Layers and I have used Drop out function as well and the activation function which I used are Relu & Sigmoid.

```
[59]: X_train = np.array(X_train)
X_test = np.array(X_test)
y_train = np.array(y_train)
y_test = np.array(y_test)
```

- Here the Input_dim is 29 as our features are also 29.

```
[60]: %%time
model_NN1 = Sequential()
model_NN1.add(layers.Dense(16,input_dim = 29,activation='relu') )
model_NN1.add(layers.Dense(units=24,activation='relu')),
model_NN1.add(Dropout(0.3))
model_NN1.add(layers.Dense(24, activation='relu'))
model_NN1.add(layers.Dense(24, activation='relu'))
model_NN1.add(layers.Dense(1, activation='sigmoid'))
model_NN1.compile(optimizer='adam',
    loss='binary_crossentropy',
    metrics=['accuracy'])
model_NN1.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 16)	480
dense_1 (Dense)	(None, 24)	408

dropout (Dropout)	(None, 24)	0
dense_2 (Dense)	(None, 24)	600
dense_3 (Dense)	(None, 24)	600
dense_4 (Dense)	(None, 1)	25

```
=====
Total params: 2,113
Trainable params: 2,113
Non-trainable params: 0
```

```
-----
CPU times: total: 938 ms
Wall time: 952 ms
```

```
[61]: %%time
model_NN1.
      ↪ compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
model_NN1.fit(X_train, y_train, batch_size=15, epochs=5)
```

```
Epoch 1/5
15190/15190 [=====] - 131s 9ms/step - loss: 0.0073 -
accuracy: 0.9991
Epoch 2/5
15190/15190 [=====] - 132s 9ms/step - loss: 0.0038 -
accuracy: 0.9993
Epoch 3/5
15190/15190 [=====] - 132s 9ms/step - loss: 0.0037 -
accuracy: 0.9993
Epoch 4/5
15190/15190 [=====] - 133s 9ms/step - loss: 0.0034 -
accuracy: 0.9993
Epoch 5/5
15190/15190 [=====] - 133s 9ms/step - loss: 0.0036 -
accuracy: 0.9993
CPU times: total: 14min 30s
Wall time: 11min 1s
```

```
[61]: <keras.callbacks.History at 0x26081293820>
```

```
[62]: score = model_NN1.evaluate(X_test, y_test)
```

```
1781/1781 [=====] - 10s 5ms/step - loss: 0.0039 -
accuracy: 0.9995
```

```
[63]: print(score)
```


[0.003913349937647581, 0.9995084404945374]

```
[64]: y_pred = model_NN1.predict(X_test)
for i in range(len(y_test)):
    if y_pred[i]>0.5:
        y_pred[i]=1
    else:
        y_pred[i]=0
cm = confusion_matrix(y_test, y_pred)
df_cm = pd.DataFrame(cm, index = (0, 1), columns = (0, 1))
plt.figure(figsize = (10,7))

sns.heatmap(df_cm, annot=True, fmt='g')
print("Test Data Accuracy: %0.4f" % accuracy_score(y_test, y_pred))
```

1781/1781 [=====] - 5s 3ms/step

Test Data Accuracy: 0.9995



```
[65]: y_pred = model_NN1.predict(X_test)
y_test = pd.DataFrame(y_test)
```

1781/1781 [=====] - 5s 3ms/step

```
[66]: cnf_matrix = confusion_matrix(y_test, y_pred.round())
```


```
[67]: print(cnf_matrix)
```

```
[[56860    8]
 [   20   74]]
```

```
[68]: cm = confusion_matrix(y_test, y_pred.round())
df_cm = pd.DataFrame(cm, index = (0, 1), columns = (0, 1))
plt.figure(figsize = (10,7))

sns.heatmap(df_cm, annot=True, fmt='g')
```

```
[68]: <AxesSubplot:>
```



```
[69]: accuracy = accuracy_score(y_test, y_pred.round())
precision = precision_score(y_test, y_pred.round())
recall = recall_score(y_test, y_pred.round())
f1 = f1_score(y_test, y_pred.round())
print('accuracy:%0.4f'%accuracy, '\tprecision:%0.4f'%precision, '\trecall:%0.
    4f'%recall, '\tF1-score:%0.4f'%f1)
```

accuracy:0.9995 precision:0.9024 recall:0.7872 F1-score:0.8409

- Here We have used multiple dense layer in the Sequential API Neural Network and we have seen a very good Accuracy Score and Precision

```
[70]: Mod_result = pd.DataFrame(['PlainNeuralNetwork', accuracy, 1-recall, recall,
    ↪precision, f1]),
    columns = ['Model', 'Accuracy', 'FalseNegRate', 'Recall',
    ↪'Precision', 'F1 Score'])
Testing_result = Testing_result.append(Mod_result, ignore_index = True)
Testing_result
```

```
[70]:
```

	Model	Accuracy	FalseNegRate	Recall	Precision	F1 Score
0	RandomForest	0.999508	0.223404	0.776596	0.912500	0.839080
1	DecisionTree	0.999421	0.244681	0.755319	0.876543	0.811429
2	Grad_Boost	0.999438	0.234043	0.765957	0.878049	0.818182
3	SGD_Classifier	0.999438	0.234043	0.765957	0.878049	0.818182
4	PlainNeuralNetwork	0.999508	0.212766	0.787234	0.902439	0.840909

- As the data is Imbalanced here i am adding Weight_class to adjust the Class Imbalance

11.1 10.1)Neural Network using Weight Class

```
[71]: %%time
y = np.array([y_train[i][0] for i in range(len(y_train))])
class_weights = class_weight.compute_class_weight(
    class_weight = "balanced",
    classes = np.unique(y),
    y = y
)
class_weights = dict(zip(np.unique(y), class_weights))
class_weights
```

CPU times: total: 109 ms

Wall time: 113 ms

```
[71]: {0: 0.5008749291043628, 1: 286.23743718592965}
```

```
[72]: %%time
model_NN1.fit(X_train,y_train,batch_size=128, epochs=10,
    ↪class_weight=class_weights, shuffle=True)
```

Epoch 1/10

1781/1781 [=====] - 15s 8ms/step - loss: 0.1513 -
accuracy: 0.9852

Epoch 2/10

1781/1781 [=====] - 15s 8ms/step - loss: 0.1024 -
accuracy: 0.9815

Epoch 3/10

```

1781/1781 [=====] - 15s 8ms/step - loss: 0.0865 -
accuracy: 0.9764
Epoch 4/10
1781/1781 [=====] - 15s 9ms/step - loss: 0.0712 -
accuracy: 0.9768
Epoch 5/10
1781/1781 [=====] - 15s 8ms/step - loss: 0.0735 -
accuracy: 0.9762
Epoch 6/10
1781/1781 [=====] - 15s 8ms/step - loss: 0.0741 -
accuracy: 0.9791
Epoch 7/10
1781/1781 [=====] - 15s 8ms/step - loss: 0.0888 -
accuracy: 0.9772
Epoch 8/10
1781/1781 [=====] - 15s 9ms/step - loss: 0.0628 -
accuracy: 0.9786
Epoch 9/10
1781/1781 [=====] - 18s 10ms/step - loss: 0.0571 -
accuracy: 0.9778
Epoch 10/10
1781/1781 [=====] - 16s 9ms/step - loss: 0.0555 -
accuracy: 0.9797
CPU times: total: 3min 52s
Wall time: 2min 34s

```

[72]: <keras.callbacks.History at 0x260d5c48820>

```

[73]: %%time
      weighted_scr = model_NN1.evaluate(X_test, y_test)

```

```

1781/1781 [=====] - 10s 5ms/step - loss: 0.0687 -
accuracy: 0.9752
CPU times: total: 13.2 s
Wall time: 9.73 s

```

```

[74]: y_pred_nn = model_NN1.predict(X_test)
      for i in range(len(y_test)):
          if y_pred_nn[i]>0.5:
              y_pred_nn[i]=1
          else:
              y_pred_nn[i]=0


      con_max = confusion_matrix(y_test, y_pred_nn)
      df_con_max = pd.DataFrame(con_max, index = (0, 1), columns = (0, 1))
      plt.figure(figsize = (10,7))

```

```
sns.heatmap(df_con_max, annot=True, fmt='g')
print("Test set Acc: %.4f" % accuracy_score(y_test, y_pred_nn))
```

1781/1781 [=====] - 5s 3ms/step

Test set Acc: 0.9752



```
[75]: accuracy = accuracy_score(y_test, y_pred_nn.round())
precision = precision_score(y_test, y_pred_nn.round())
recall = recall_score(y_test, y_pred_nn.round())
f1 = f1_score(y_test, y_pred_nn.round())

print('accuracy: %.4f' % accuracy, '\tprecision: %.4f' % precision, '\trecall: %.4f' % recall, '\tF1-score: %.4f' % f1)
```

accuracy:0.9752 precision:0.0572 recall:0.9043 F1-score:0.1076

- The thing here to notice is that we have very much improved in Detecting the **Fraudulent Transactions** but our **Accuracy** is decreased we have very low rate of False Negative which is most important criteria to classify the **Fraudulent Transaction**. On other hand we have seen that the false positive score is increased which means it will misclassify the normal transactions.


```
[76]: Mod_result = pd.DataFrame(['WeightedNeuralNetwork', accuracy, 1-recall,
    ↪recall, precision, f1]),
    columns = ['Model', 'Accuracy', 'FalseNegRate', 'Recall',
    ↪'Precision', 'F1 Score'])

Testing_result = Testing_result.append(Mod_result, ignore_index = True)
Testing_result
```

```
[76]:
```

	Model	Accuracy	FalseNegRate	Recall	Precision	\
0	RandomForest	0.999508	0.223404	0.776596	0.912500	
1	DecisionTree	0.999421	0.244681	0.755319	0.876543	
2	Grad_Boost	0.999438	0.234043	0.765957	0.878049	
3	SGD_Classifier	0.999438	0.234043	0.765957	0.878049	
4	PlainNeuralNetwork	0.999508	0.212766	0.787234	0.902439	
5	WeightedNeuralNetwork	0.975247	0.095745	0.904255	0.057201	


```

F1 Score
0 0.839080
1 0.811429
2 0.818182
3 0.818182
4 0.840909
5 0.107595
```

11.2 10.2) Neural Network using Undersampling Technique

This is one more approach to balance the train set , using this undersampling technique will help us to randomly pick the Normal and fraudulent transactions from the data set.

```
[77]: fraud_transactions = np.array(CC_data[CC_data.Class == 1].index)
    fraud_trans_numbers = len(fraud_transactions)
    print(fraud_trans_numbers)
```

492

```
[78]: non_fraudlent_trans = CC_data[CC_data.Class == 0].index

    len(non_fraudlent_trans)
```

[78]: 284315

```
[79]: Ran_non_fraudlent_trans = np.random.choice(non_fraudlent_trans,
    ↪fraud_trans_numbers, replace=False)
    Ran_non_fraudlent_trans = np.array(Ran_non_fraudlent_trans)
    print(len(Ran_non_fraudlent_trans))
```

492

```
[80]: trans_underSample = np.concatenate([fraud_transactions,Ran_non_fraudlent_trans])
      print(len(trans_underSample))
```

984

```
[81]: df_underSample = CC_data.iloc[trans_underSample,:]
```

```
[82]: X_under_Sample = df_underSample.iloc[:,df_underSample.columns != 'Class']
      y_under_Sample = df_underSample.iloc[:,df_underSample.columns == 'Class']
```

```
[83]: X_train, X_test, y_train, y_test = \
      ↪train_test_split(X_under_Sample,y_under_Sample, test_size=0.3)
```

```
[84]: X_train = np.array(X_train)
      X_test = np.array(X_test)
      y_train = np.array(y_train)
      y_test = np.array(y_test)
```

```
[85]: X_train.shape
```

```
[85]: (688, 29)
```

```
[86]: X_test.shape
```

```
[86]: (296, 29)
```

```
[87]: model_NN1.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 16)	480
dense_1 (Dense)	(None, 24)	408
dropout (Dropout)	(None, 24)	0
dense_2 (Dense)	(None, 24)	600
dense_3 (Dense)	(None, 24)	600
dense_4 (Dense)	(None, 1)	25

Total params: 2,113
Trainable params: 2,113

Non-trainable params: 0

- I am using the same neural network again after undersampling

```
[88]: %%time
model_NN1.
      ↪ compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

model_NN1.fit(X_train, y_train, batch_size=128, epochs=10)
```

```
Epoch 1/10
6/6 [=====] - 1s 11ms/step - loss: 0.0943 - accuracy: 0.9724
Epoch 2/10
6/6 [=====] - 0s 12ms/step - loss: 0.0855 - accuracy: 0.9826
Epoch 3/10
6/6 [=====] - 0s 12ms/step - loss: 0.0885 - accuracy: 0.9811
Epoch 4/10
6/6 [=====] - 0s 12ms/step - loss: 0.0803 - accuracy: 0.9767
Epoch 5/10
6/6 [=====] - 0s 12ms/step - loss: 0.0679 - accuracy: 0.9826
Epoch 6/10
6/6 [=====] - 0s 11ms/step - loss: 0.0677 - accuracy: 0.9840
Epoch 7/10
6/6 [=====] - 0s 11ms/step - loss: 0.0646 - accuracy: 0.9840
Epoch 8/10
6/6 [=====] - 0s 11ms/step - loss: 0.0693 - accuracy: 0.9811
Epoch 9/10
6/6 [=====] - 0s 11ms/step - loss: 0.0592 - accuracy: 0.9840
Epoch 10/10
6/6 [=====] - 0s 9ms/step - loss: 0.0640 - accuracy: 0.9840
CPU times: total: 1.59 s
Wall time: 1.3 s
```

```
[88]: <keras.callbacks.History at 0x260ade6dee0>
```

```
[89]: y_pred_u = model_NN1.predict(X_test)
      for i in range(len(y_test)):
```

```

if y_pred_u[i]>0.5:
    y_pred_u[i]=1
else:
    y_pred_u[i]=0


con_max = confusion_matrix(y_test, y_pred_u) # rows = truth, cols = prediction
df_con_max = pd.DataFrame(con_max, index = (0, 1), columns = (0, 1))
plt.figure(figsize = (10,7))

sns.heatmap(df_con_max, annot=True, fmt='g')
print("Test set Acc: %.4f" % accuracy_score(y_test, y_pred_u))

```

10/10 [=====] - 0s 4ms/step

Test set Acc: 0.9595



- This model seems to be pretty accurate on test set.

```

[90]: accuracy = accuracy_score(y_test, y_pred_u.round())
precision = precision_score(y_test, y_pred_u.round())
recall = recall_score(y_test, y_pred_u.round())

```

```
f1 = f1_score(y_test, y_pred_u.round())

print('accuracy:%0.4f'%accuracy, '\tprecision:%0.4f'%precision, '\trecall:%0.4f'%recall, '\tF1-score:%0.4f'%f1)
```

```
accuracy:0.9595          precision:0.9735          recall:0.9484    F1-score:0.9608
```

```
[91]: Mod_result = pd.DataFrame(['UnderSampledNeuralNetwork', accuracy, 1-recall, recall, precision, f1]),
        columns = ['Model', 'Accuracy', 'FalseNegRate', 'Recall', 'Precision', 'F1 Score'])

Testing_result = Testing_result.append(Mod_result, ignore_index = True)
Testing_result
```

```
[91]:
```

	Model	Accuracy	FalseNegRate	Recall	Precision	\
0	RandomForest	0.999508	0.223404	0.776596	0.912500	
1	DecisionTree	0.999421	0.244681	0.755319	0.876543	
2	Grad_Boost	0.999438	0.234043	0.765957	0.878049	
3	SGD_Classifier	0.999438	0.234043	0.765957	0.878049	
4	PlainNeuralNetwork	0.999508	0.212766	0.787234	0.902439	
5	WeightedNeuralNetwork	0.975247	0.095745	0.904255	0.057201	
6	UnderSampledNeuralNetwork	0.959459	0.051613	0.948387	0.973510	

	F1 Score
0	0.839080
1	0.811429
2	0.818182
3	0.818182
4	0.840909
5	0.107595
6	0.960784

11.3 10.3) Neural Network using SMOTE

- As I have use undersampling before SMOTE (Synthetic Minority Oversample Technique) is opposite of undersampling it upsample the minority class to same level of Majority Class
- The SMOTE generates a new Vector between two 2 existing class. This will increase the Fraudulent class.

```
[92]: X = CC_data.iloc[:, CC_data.columns != 'Class']
      y = CC_data.iloc[:, CC_data.columns == 'Class']
```

```
[93]: print(X.shape)
      print(y.shape)
```

```
if X.shape[0] != y.shape[0]:
    print("X and y rows are mismatched, check dataset again")
```

```
(284807, 29)
```

```
(284807, 1)
```

```
[94]: oversample = SMOTE()
      X_resam, y_resam = oversample.fit_resample(X, y.values.ravel())
```

```
[95]: y_resam
```

```
[95]: array([0, 0, 0, ..., 1, 1, 1], dtype=int64)
```

```
[96]: print('Total no of transactions before applying SMOTE: ', len(y), '...after_
      ↪SMOTE : ', len(y_resam))
      print('Total no of transactions before applying SMOTE: ', len(y[y.Class==1]),
            '...after SMOTE : ', np.sum(y_resam[y_resam==1]))
```

```
Total no of transactions before applying SMOTE: 284807 ...after SMOTE : 568630
```

```
Total no of transactions before applying SMOTE: 492 ...after SMOTE : 284315
```

```
[97]: y_resamp = pd.DataFrame(y_resam)
      X_resamp = pd.DataFrame(X_resam)
```

```
[98]: X_train, X_test, y_train, y_test =_
      ↪train_test_split(X_resamp,y_resamp,test_size=0.2)
```

```
[99]: X_train = np.array(X_train)
      X_test = np.array(X_test)
      y_train = np.array(y_train)
      y_test = np.array(y_test)
```

```
[100]: %%time
model_NN2 = Sequential([
    Dense(units=16, input_dim = 29,activation='relu'),
    Dense(units=24,activation='relu'),
    Dropout(0.5),
    Dense(24,activation='relu'),
    Dense(24,activation='relu'),
    Dense(1,activation='sigmoid'),
])
```

```
CPU times: total: 46.9 ms
```

```
Wall time: 47.9 ms
```

```
[101]: %%time
```


```
model_NN2.  
    ↪ compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])  
model_NN2.fit(X_train, y_train, batch_size= 128, epochs=5 )
```

```
Epoch 1/5  
3554/3554 [=====] - 33s 9ms/step - loss: 0.0754 -  
accuracy: 0.9686  
Epoch 2/5  
3554/3554 [=====] - 29s 8ms/step - loss: 0.0256 -  
accuracy: 0.9918  
Epoch 3/5  
3554/3554 [=====] - 30s 8ms/step - loss: 0.0187 -  
accuracy: 0.9946  
Epoch 4/5  
3554/3554 [=====] - 30s 8ms/step - loss: 0.0150 -  
accuracy: 0.9959  
Epoch 5/5  
3554/3554 [=====] - 30s 8ms/step - loss: 0.0128 -  
accuracy: 0.9966  
CPU times: total: 3min 23s  
Wall time: 2min 31s
```

[101]: <keras.callbacks.History at 0x260d5c20040>

```
[102]: y_predict = model_NN2.predict(X_test)  
y_ex = pd.DataFrame(y_test)  
  
for i in range(len(y_ex)):  
    if y_predict[i]>0.5:  
        y_predict[i]=1  
    else:  
        y_predict[i]=0  
  
con_max = confusion_matrix(y_ex, y_predict) # rows = truth, cols = prediction  
df_con_max = pd.DataFrame(con_max, index = (0, 1), columns = (0, 1))  
plt.figure(figsize = (10,7))  
  
sns.heatmap(df_con_max, annot=True, fmt='g')  
print("Test set Acc: %0.4f" % accuracy_score(y_ex, y_predict))
```

```
3554/3554 [=====] - 11s 3ms/step  
Test set Acc: 0.9979
```



```
[103]: accuracy = accuracy_score(y_test, y_predict.round())
precision = precision_score(y_test, y_predict.round())
recall = recall_score(y_test, y_predict.round())
f1 = f1_score(y_test, y_predict.round())

print('accuracy:%0.4f'%accuracy, '\tprecision:%0.4f'%precision, '\trecall:%0.
↪4f'%recall, '\tF1-score:%0.4f'%f1)
```

accuracy:0.9979 precision:0.9966 recall:0.9992 F1-score:0.9979

11.3.1 10.3.1) Testing on Full Data Set Now

```
[104]: y_predict = model_NN2.predict(X)
y_ex = pd.DataFrame(y)

for i in range(len(y_ex)):
    if y_predict[i]>0.5:
        y_predict[i]=1
    else:
        y_predict[i]=0


con_max = confusion_matrix(y_ex, y_predict) # rows = truth, cols = prediction
```



```
df_con_max = pd.DataFrame(con_max, index = (0, 1), columns = (0, 1))
plt.figure(figsize = (10,7))

sns.heatmap(df_con_max, annot=True, fmt='g')
print("Test set Acc: %0.4f" % accuracy_score(y_ex, y_predict))
```

8901/8901 [=====] - 25s 3ms/step
Test set Acc: 0.9970



Now we can say that after using SMOTE it's clear that our model is predicting very precisely it predicted the 487 out 492 to be fraudulent which is a very good score.

12 11) Model Evaluation

```
[105]: Mod_result = pd.DataFrame([['OverSampledNeuralNetwork', accuracy, 1-recall,
    ↪recall, precision, f1]],
    columns = ['Model', 'Accuracy', 'FalseNegRate', 'Recall',
    ↪'Precision', 'F1 Score'])

Testing_result = Testing_result.append(Mod_result, ignore_index = True)
```

Testing_result

```
[105]:
```

	Model	Accuracy	FalseNegRate	Recall	Precision	\
0	RandomForest	0.999508	0.223404	0.776596	0.912500	
1	DecisionTree	0.999421	0.244681	0.755319	0.876543	
2	Grad_Boost	0.999438	0.234043	0.765957	0.878049	
3	SGD_Classifier	0.999438	0.234043	0.765957	0.878049	
4	PlainNeuralNetwork	0.999508	0.212766	0.787234	0.902439	
5	WeightedNeuralNetwork	0.975247	0.095745	0.904255	0.057201	
6	UnderSampledNeuralNetwork	0.959459	0.051613	0.948387	0.973510	
7	OverSampledNeuralNetwork	0.997898	0.000791	0.999209	0.996597	

	F1 Score
0	0.839080
1	0.811429
2	0.818182
3	0.818182
4	0.840909
5	0.107595
6	0.960784
7	0.997901

13 12) Conclusion and Actionable Insights

While working on this Anomaly detection task which was related to **Credit Card Transaction** we have started with **Exploratory Data Analysis** and we found out that given data was highly Imbalanced, later on I tried to find the correlations between data and found out that **Time** feature was not relevant in **Fraud Detection** so we did some data preparation and done some data scaling later on after scaling I have used multiple classification algorithms and even used **Random Grid Search** and **Ensemble Learning** and then we moved towards **Neural Network** while using that I tried to use both **Over Sampling** and **Undersampling** and **Weight class** technique to fee the data in **Neural Network** so after all of that I have seen that using **SMOTE (Oversampling)** we got the desired results as it was able to caught the most precised amount of **Fraudulent Transaction**

What we can do to improve this pipeline is that the data is very much Imbalance we need more of **Fraudulent Data** to make over pipeline more accurate and we can run some pretrained model as well on this dataset it will may increase the desired result.